

# Agentic Efficiency Report

agentic-crm · zip import · scan SCN-2026-01-01-0001

75

Watchlist + PartialScan

Weighted efficiency verdict

## What we found

|                                |                         |
|--------------------------------|-------------------------|
| Overall efficiency score       | 75 / 100                |
| Verdict                        | Watchlist (PartialScan) |
| Files discovered               | 94                      |
| Files parsed                   | 69                      |
| Files skipped                  | 25                      |
| Tokens analysed                | 55,038                  |
| Estimated monthly token waste  | 14,711,970              |
| Estimated monthly credit waste | 147.12 credits          |
| Estimated monthly dollar waste | \$82.38                 |
| Savings range (low / high)     | \$65.90 - \$98.86       |
| Duplicate clusters found       | 2                       |
| Oversized context files        | 3                       |
| Overlapping agent groups       | 1                       |
| Review stages inferred         | 10                      |
| Agent-like files detected      | 46                      |

## Category scores

| Category                  | Score  | Weight | Penalty | What this means                                                                |
|---------------------------|--------|--------|---------|--------------------------------------------------------------------------------|
| Redundancy                | 52/100 | 25%    | -48     | 2 groups of near-identical files and 4 repeated instruction blocks were found. |
| Token bloat               | 82/100 | 25%    | -18     | 3 context or memory files are over 8,000 tokens.                               |
| Review overhead           | 85/100 | 20%    | -15     | 10 review or approval steps were inferred from your instructions.              |
| Agent sprawl              | 80/100 | 20%    | -20     | 46 agent-style files were found (20 agent roles, 13 skills).                   |
| Architecture inefficiency | 88/100 | 10%    | -12     | 9 service directories against 9 hand written source files.                     |

## Top 5 waste drivers, in plain language

### 1. The same instruction block is pasted into 3 files

One block of instructions has been copied into several agent files. Only one copy is needed. The rest is paid for on every run.

Redundancy · 6,268,000 tokens/month · 62.68 credits · \$35.10

### 2. Context file is very large: memory.md

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

Token bloat · 2,535,200 tokens/month · 25.35 credits · \$14.20

### 3. Context file is very large: context.md

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

Token bloat · 1,553,200 tokens/month · 15.53 credits · \$8.70

4. 20 files are near copies of each other

These files say almost the same thing. Every time your agents run, the same words get sent to the model more than once. Keeping one source of truth removes the repeat cost.

Redundancy · 1,055,200 tokens/month · 10.55 credits · \$5.91

5. 46 agent-style files is more than this repo needs

There are many separate agent, skill and prompt files. Each one adds instructions that have to be loaded and kept in sync. Fewer, clearer files cost less and break less often.

Agent sprawl · 816,000 tokens/month · 8.16 credits · \$4.57

Savings assumptions and formulas

- `estimated_tokens = ceil(character_count / 4)`
- Report credits: 1 credit = 100,000 tokens.
- Vendor billing: \$1.00 = 100 vendor credits; 1 vendor credit = 2,500 input tokens or 500 output tokens (output costs 5x input).
- Derived rates: \$4.00 per 1M input tokens, \$20.00 per 1M output tokens.
- Waste is assumed to be 90% input tokens and 10% output tokens.
- Each agent asset is assumed to be loaded 200 times per month.
- Aggregate savings show a +/-20% range.
- Similarity uses normalised text so whitespace-only and case-only changes are ignored.
- Files under 200 characters are excluded from duplicate detection.

| Setting                         | Value                                               |
|---------------------------------|-----------------------------------------------------|
| Tokens per report credit        | 100,000                                             |
| Vendor credits per dollar       | 100.0                                               |
| Input tokens per vendor credit  | 2,500                                               |
| Output tokens per vendor credit | 500                                                 |
| \$ per 1M input tokens          | \$4.00                                              |
| \$ per 1M output tokens         | \$20.00                                             |
| Agent runs per month            | 200                                                 |
| Output token share of waste     | 10%                                                 |
| Aggregate variance              | +/-20%                                              |
| Rates last refreshed            | never                                               |
| Rates source                    | Product owner supplied defaults (not yet refreshed) |

How the score was calculated

| Penalty rule                                                     | Hits | Points each | Cap | Applied | Category                  |
|------------------------------------------------------------------|------|-------------|-----|---------|---------------------------|
| Near-duplicate agent/context file pair (similarity >= 0.80)      | 20   | 5           | 32  | -32     | redundancy                |
| Repeated instruction block >= 150 chars in 3+ files              | 4    | 4           | 20  | -16     | redundancy                |
| Context or memory file over 8,000 tokens                         | 3    | 6           | 24  | -18     | token_bloat               |
| Agent-like files above 12                                        | 34   | 2           | 20  | -20     | agent_sprawl              |
| Review/approval stages above 4                                   | 6    | 5           | 15  | -15     | review_overhead           |
| Microservice mismatch (>= 8 service dirs and < 120 source files) | 1    | 12          | 12  | -12     | architecture_inefficiency |
| Monolith mismatch (1 service and >= 10 overlapping agent roles)  | 0    | 10          | 10  | -0      | architecture_inefficiency |

## Findings (14 issues)

| Issue               | Severity | Category                  | Title                                                     | Impacted files                                                                                                    | Tokens /mo | Credits/mo | \$/mo   |
|---------------------|----------|---------------------------|-----------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|------------|------------|---------|
| ISS-2026-07-29-0006 | critical | redundancy                | The same instruction block is pasted into 3 files         | context/context.md, context/product-context.md, memory/memory.md                                                  | 6,268,000  | 62.68      | \$35.10 |
| ISS-2026-07-29-0007 | high     | token_bloat               | Context file is very large: memory.md                     | memory/memory.md                                                                                                  | 2,535,200  | 25.35      | \$14.20 |
| ISS-2026-07-29-0008 | high     | token_bloat               | Context file is very large: context.md                    | context/context.md                                                                                                | 1,553,200  | 15.53      | \$8.70  |
| ISS-2026-07-29-0001 | high     | redundancy                | 20 files are near copies of each other                    | agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md (+16 more) | 1,055,200  | 10.55      | \$5.91  |
| ISS-2026-07-29-0010 | medium   | agent_sprawl              | 46 agent-style files is more than this repo needs         | agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md (+21 more) | 816,000    | 8.16       | \$4.57  |
| ISS-2026-07-29-0003 | medium   | redundancy                | The same instruction block is pasted into 43 files        | agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md (+39 more) | 764,400    | 7.64       | \$4.28  |
| ISS-2026-07-29-0012 | high     | architecture_inefficiency | Split into many services but there is little code in them | services/svc-1, services/svc-2, services/svc-3, services/svc-4 (+5 more)                                          | 546,670    | 5.47       | \$3.06  |
| ISS-2026-07-29-0009 | medium   | token_bloat               | Context file is very large: product-context.md            | context/product-context.md                                                                                        | 401,000    | 4.01       | \$2.25  |
| ISS-2026-07-29-0004 | medium   | redundancy                | The same instruction block is pasted into 23 files        | agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md (+19 more) | 391,600    | 3.92       | \$2.19  |
| ISS-2026-07-29-0005 | low      | redundancy                | The same instruction block is pasted into 20 files        | agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md (+16 more) | 182,400    | 1.82       | \$1.02  |
| ISS-2026-07-29-0013 | low      | agent_sprawl              | 13 skill files, and some repeat each other                | skills/archive.skill.md, skills/classify.skill.md, skills/escalate.skill.md, skills/extract.skill.md (+9 more)    | 72,000     | 0.72       | \$0.40  |
| ISS-2026-07-29-0011 | low      | review_overhead           | 10 separate review or approval steps were found           | agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md            | 68,100     | 0.68       | \$0.38  |
| ISS-2026-07-29-0002 | low      | redundancy                | 2 files are near copies of each other                     | instructions.md, orchestrator.md                                                                                  | 38,200     | 0.38       | \$0.21  |
| ISS-2026-07-29-0014 | low      | review_overhead           | 10 orchestration files create extra hand-offs             | instruction.md, instructions.md, orchestrator.md, prompts/stage-1.prompt.md (+6 more)                             | 20,000     | 0.20       | \$0.11  |

## Evidence detail

### ***ISS-2026-07-29-0006 — The same instruction block is pasted into 3 files***

One block of instructions has been copied into several agent files. Only one copy is needed. The rest is paid for on every run.

**Evidence:** Block of 62678 characters (~15670 tokens) found in: context/context.md, context/product-context.md, memory/memory.md | Sample: "The customer has an established relationship with the platform and has raised several tickets about billing, onboarding and reporting over the last four quarter"

**Formula:**  $\text{monthly\_waste} = \text{block\_tokens} \times (\text{files\_containing\_block} - 1) \times \text{runs\_per\_month}$

**Recommendation:** Move the shared block into one file and reference it from the others. (impact Medium, effort Small)

### ***ISS-2026-07-29-0007 — Context file is very large: memory.md***

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

**Evidence:** memory/memory.md is about 20,676 tokens which is 12,676 tokens over the 8,000 token guideline.

**Formula:**  $\text{monthly\_waste} = (\text{file\_tokens} - 8000) \times \text{runs\_per\_month}$

**Recommendation:** Split this file into a short always-on summary plus detail files that load only when needed. (impact Medium, effort Medium)

### ***ISS-2026-07-29-0008 — Context file is very large: context.md***

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

**Evidence:** context/context.md is about 15,766 tokens which is 7,766 tokens over the 8,000 token guideline.

**Formula:**  $\text{monthly\_waste} = (\text{file\_tokens} - 8000) \times \text{runs\_per\_month}$

**Recommendation:** Split this file into a short always-on summary plus detail files that load only when needed. (impact Medium, effort Medium)

### ***ISS-2026-07-29-0001 — 20 files are near copies of each other***

These files say almost the same thing. Every time your agents run, the same words get sent to the model more than once. Keeping one source of truth removes the repeat cost.

**Evidence:** Group DUP-001 - highest similarity 100%. Files: agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md, agents/intake.agent.md, agents/legacy/intake-backup.agent.md ...

**Formula:**  $\text{monthly\_waste} = (\text{sum of duplicate file tokens} - \text{largest file tokens}) \times \text{runs\_per\_month}$

**Recommendation:** Keep the most complete file, delete or shrink the copies, and link to the kept file instead. (impact High, effort Medium)

### ***ISS-2026-07-29-0010 — 46 agent-style files is more than this repo needs***

There are many separate agent, skill and prompt files. Each one adds instructions that have to be loaded and kept in sync. Fewer, clearer files cost less and break less often.

**Evidence:** 46 agent-like files detected against a healthy guideline of 12. Median agent file size is 240 tokens.

**Formula:**  $\text{monthly\_waste} = \text{extra\_agent\_files} \times \text{median\_agent\_tokens} \times \text{runs\_per\_month} \times 0.5$

**Recommendation:** Group agents by job. Merge the ones that overlap and delete the ones nobody calls. (impact High, effort Large)

### ***ISS-2026-07-29-0003 — The same instruction block is pasted into 43 files***

One block of instructions has been copied into several agent files. Only one copy is needed. The rest is paid for on every run.

**Evidence:** Block of 363 characters (~91 tokens) found in: agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md, agents/intake.agent.md, agents/legacy/intake-backup.agent.md ... | Sample: "Before you begin any task you must read the repository conventions file, confirm the ticket number, check that the branch name follows the release naming policy"

**Formula:**  $\text{monthly\_waste} = \text{block\_tokens} \times (\text{files\_containing\_block} - 1) \times \text{runs\_per\_month}$

**Recommendation:** Move the shared block into one file and reference it from the others. (impact High, effort Small)

### ***ISS-2026-07-29-0012 — Split into many services but there is little code in them***

The repo is arranged as many small services, yet the amount of real code is small. Each service still needs its own setup and its own instructions, so the overhead is larger than the work being done.

**Evidence:** 9 service directories detected with only 9 hand written source files.

**Formula:**  $\text{monthly\_waste} = 5\% \times \text{total\_agent\_asset\_tokens} \times \text{runs\_per\_month}$

**Recommendation:** Fold the thin services back together until each one has a clear, separate job. (impact High, effort Large)

### ***ISS-2026-07-29-0009 — Context file is very large: product-context.md***

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

**Evidence:** context/product-context.md is about 10,005 tokens which is 2,005 tokens over the 8,000 token guideline.

**Formula:**  $\text{monthly\_waste} = (\text{file\_tokens} - 8000) \times \text{runs\_per\_month}$

**Recommendation:** Split this file into a short always-on summary plus detail files that load only when needed. (impact Medium, effort Medium)

### ISS-2026-07-29-0004 — The same instruction block is pasted into 23 files

One block of instructions has been copied into several agent files. Only one copy is needed. The rest is paid for on every run.

**Evidence:** Block of 353 characters (~89 tokens) found in: agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md, agents/intake.agent.md, agents/legacy/intake-backup.agent.md ... | Sample: "Every change moves through the following gates in order: self review, then peer review, then code review by the owning team, then a security review, then archit"

**Formula:**  $\text{monthly\_waste} = \text{block\_tokens} \times (\text{files\_containing\_block} - 1) \times \text{runs\_per\_month}$

**Recommendation:** Move the shared block into one file and reference it from the others. (impact High, effort Small)

### ISS-2026-07-29-0005 — The same instruction block is pasted into 20 files

One block of instructions has been copied into several agent files. Only one copy is needed. The rest is paid for on every run.

**Evidence:** Block of 190 characters (~48 tokens) found in: agents/billing.agent.md, agents/escalation.agent.md, agents/forecast.agent.md, agents/handoff.agent.md, agents/intake.agent.md, agents/legacy/intake-backup.agent.md ... | Sample: "When finished, write a summary into the shared memory file and notify the orchestrator so the next stage can start. Include the ticket number, the files touched"

**Formula:**  $\text{monthly\_waste} = \text{block\_tokens} \times (\text{files\_containing\_block} - 1) \times \text{runs\_per\_month}$

**Recommendation:** Move the shared block into one file and reference it from the others. (impact High, effort Small)

### ISS-2026-07-29-0013 — 13 skill files, and some repeat each other

Skill files are meant to be small and specific. When there are many of them and they repeat, the model reads the same guidance several times.

**Evidence:** 13 skill files detected; 0 of them sit inside near-duplicate groups.

**Formula:**  $\text{monthly\_waste} = \text{duplicated\_skill\_tokens} \times \text{runs\_per\_month} \times 0.5$

**Recommendation:** Keep one skill file per capability and delete the near copies. (impact Medium, effort Small)

### ISS-2026-07-29-0011 — 10 separate review or approval steps were found

Your instructions describe many checks before work is accepted. Each extra check means another pass over the same material, which costs time and money.

**Evidence:** Stages inferred: Architecture review, Code review, Escalation review, Final approval, Human in the loop, Peer review, Product approval, Security review, Self review, Sign-off. Examples: Architecture review -> agents/billing.agent.md; Code review -> agents/billing.agent.md; Escalation review -> agents/billing.agent.md; Final approval -> agents/billing.agent.md

**Formula:**  $\text{monthly\_waste} = \text{extra\_review\_stages} \times \text{avg\_instruction\_tokens} \times \text{runs\_per\_month} \times 0.25$

**Recommendation:** Keep at most four review gates. Combine the overlapping ones into a single checklist. (impact High, effort Medium)

### ISS-2026-07-29-0002 — 2 files are near copies of each other

These files say almost the same thing. Every time your agents run, the same words get sent to the model more than once. Keeping one source of truth removes the repeat cost.

**Evidence:** Group DUP-002 - highest similarity 84%. Files: instructions.md, orchestrator.md

**Formula:**  $\text{monthly\_waste} = (\text{sum of duplicate file tokens} - \text{largest file tokens}) \times \text{runs\_per\_month}$

**Recommendation:** Keep the most complete file, delete or shrink the copies, and link to the kept file instead. (impact Medium, effort Small)

### ISS-2026-07-29-0014 — 10 orchestration files create extra hand-offs

There are many files describing how work moves between agents. Every hand-off adds instructions that must be loaded and kept in step with the others.

**Evidence:** 10 orchestration or prompt files detected, average size 125 tokens.

**Formula:**  $\text{monthly\_waste} = (\text{orchestration\_files} - 6) \times \text{avg\_orchestration\_tokens} \times \text{runs\_per\_month} \times 0.2$

**Recommendation:** Describe the flow once in a single orchestrator file and let the agents stay simple. (impact Medium, effort Medium)

## Recommended actions, ranked by impact then effort

| # | Action                                                                                       | Category        | Impact | Effort | Tokens/mo saved | \$/mo saved |
|---|----------------------------------------------------------------------------------------------|-----------------|--------|--------|-----------------|-------------|
| 1 | Move the shared block into one file and reference it from the others.                        | Redundancy      | High   | Small  | 764,400         | \$4.28      |
| 2 | Keep the most complete file, delete or shrink the copies, and link to the kept file instead. | Redundancy      | High   | Medium | 1,055,200       | \$5.91      |
| 3 | Keep at most four review gates. Combine the overlapping ones into a single checklist.        | Review overhead | High   | Medium | 68,100          | \$0.38      |
| 4 | Group agents by job. Merge the ones that overlap and delete the ones nobody calls.           | Agent sprawl    | High   | Large  | 816,000         | \$4.57      |

| # | Action                                                                                       | Category                  | Impact | Effort | Tokens/mo saved | \$/mo saved |
|---|----------------------------------------------------------------------------------------------|---------------------------|--------|--------|-----------------|-------------|
| 5 | Fold the thin services back together until each one has a clear, separate job.               | Architecture inefficiency | High   | Large  | 546,670         | \$3.06      |
| 6 | Keep one skill file per capability and delete the near copies.                               | Agent sprawl              | Medium | Small  | 72,000          | \$0.40      |
| 7 | Split this file into a short always-on summary plus detail files that load only when needed. | Token bloat               | Medium | Medium | 2,535,200       | \$14.20     |
| 8 | Describe the flow once in a single orchestrator file and let the agents stay simple.         | Review overhead           | Medium | Medium | 20,000          | \$0.11      |

File inventory by category

| Group                    | Files | Parsed | Skipped | Tokens |
|--------------------------|-------|--------|---------|--------|
| Agents                   | 20    | 20     | 0       | 5,564  |
| Skills                   | 13    | 13     | 0       | 1,403  |
| Context and memory       | 3     | 3      | 0       | 46,447 |
| Prompt and orchestration | 10    | 10     | 0       | 1,253  |
| Source code              | 9     | 9      | 0       | 90     |
| Diagrams                 | 1     | 1      | 0       | 13     |
| Docs                     | 2     | 2      | 0       | 207    |
| Other text assets        | 36    | 11     | 25      | 61     |

Largest 94 files by estimated tokens

| Path                                 | Category       | Status  | Lines | Size (KB) | Tokens | Dup group |
|--------------------------------------|----------------|---------|-------|-----------|--------|-----------|
| memory/memory.md                     | context_memory | Scanned | 8     | 80.8      | 20,676 | -         |
| context/context.md                   | context_memory | Scanned | 7     | 61.6      | 15,766 | -         |
| context/product-context.md           | context_memory | Scanned | 5     | 39.1      | 10,005 | -         |
| agents/legacy/intake-backup.agent.md | agent          | Scanned | 23    | 1.1       | 288    | DUP-001   |
| agents/intake.agent.md               | agent          | Scanned | 23    | 1.1       | 280    | DUP-001   |
| agents/legacy/intake-copy.agent.md   | agent          | Scanned | 23    | 1.1       | 280    | DUP-001   |
| agents/legacy/intake-old.agent.md    | agent          | Scanned | 23    | 1.1       | 280    | DUP-001   |
| agents/escalation.agent.md           | agent          | Scanned | 23    | 1.1       | 279    | DUP-001   |
| agents/legacy/intake-v2.agent.md     | agent          | Scanned | 23    | 1.1       | 279    | DUP-001   |
| agents/onboarding.agent.md           | agent          | Scanned | 23    | 1.1       | 279    | DUP-001   |
| agents/forecast.agent.md             | agent          | Scanned | 23    | 1.1       | 278    | DUP-001   |
| agents/proposal.agent.md             | agent          | Scanned | 23    | 1.1       | 278    | DUP-001   |
| agents/reporting.agent.md            | agent          | Scanned | 23    | 1.1       | 278    | DUP-001   |
| agents/billing.agent.md              | agent          | Scanned | 23    | 1.1       | 277    | DUP-001   |
| agents/handoff.agent.md              | agent          | Scanned | 23    | 1.1       | 277    | DUP-001   |
| agents/notes.agent.md                | agent          | Scanned | 23    | 1.1       | 277    | DUP-001   |
| agents/renewal.agent.md              | agent          | Scanned | 23    | 1.1       | 277    | DUP-001   |
| agents/research.agent.md             | agent          | Scanned | 23    | 1.1       | 277    | DUP-001   |
| agents/support.agent.md              | agent          | Scanned | 23    | 1.1       | 277    | DUP-001   |
| agents/legal.agent.md                | agent          | Scanned | 23    | 1.1       | 276    | DUP-001   |
| agents/pricing.agent.md              | agent          | Scanned | 23    | 1.1       | 276    | DUP-001   |
| agents/triage.agent.md               | agent          | Scanned | 23    | 1.1       | 276    | DUP-001   |
| agents/qa.agent.md                   | agent          | Scanned | 23    | 1.1       | 275    | DUP-001   |
| orchestrator.md                      | orchestration  | Scanned | 15    | 0.8       | 206    | DUP-002   |
| instructions.md                      | orchestration  | Scanned | 7     | 0.7       | 191    | DUP-002   |
| prompts/stage-1.prompt.md            | orchestration  | Scanned | 5     | 0.4       | 109    | -         |
| prompts/stage-2.prompt.md            | orchestration  | Scanned | 5     | 0.4       | 109    | -         |
| prompts/stage-3.prompt.md            | orchestration  | Scanned | 5     | 0.4       | 109    | -         |
| prompts/stage-4.prompt.md            | orchestration  | Scanned | 5     | 0.4       | 109    | -         |
| prompts/stage-5.prompt.md            | orchestration  | Scanned | 5     | 0.4       | 109    | -         |
| prompts/stage-6.prompt.md            | orchestration  | Scanned | 5     | 0.4       | 109    | -         |
| skills/reconcile.skill.md            | skill          | Scanned | 5     | 0.4       | 109    | -         |
| skills/summarise.skill.md            | skill          | Scanned | 5     | 0.4       | 109    | -         |
| skills/translate.skill.md            | skill          | Scanned | 5     | 0.4       | 109    | -         |
| skills/archive.skill.md              | skill          | Scanned | 5     | 0.4       | 108    | -         |
| skills/classify.skill.md             | skill          | Scanned | 5     | 0.4       | 108    | -         |
| skills/escalate.skill.md             | skill          | Scanned | 5     | 0.4       | 108    | -         |

| Path                              | Category      | Status             | Lines | Size (KB) | Tokens | Dup group |
|-----------------------------------|---------------|--------------------|-------|-----------|--------|-----------|
| skills/extract.skill.md           | skill         | Scanned            | 5     | 0.4       | 108    | -         |
| skills/schedule.skill.md          | skill         | Scanned            | 5     | 0.4       | 108    | -         |
| skills/validate.skill.md          | skill         | Scanned            | 5     | 0.4       | 108    | -         |
| skills/format.skill.md            | skill         | Scanned            | 5     | 0.4       | 107    | -         |
| skills/legacy/search-old.skill.md | skill         | Scanned            | 5     | 0.4       | 107    | -         |
| skills/notify.skill.md            | skill         | Scanned            | 5     | 0.4       | 107    | -         |
| skills/search.skill.md            | skill         | Scanned            | 5     | 0.4       | 107    | -         |
| instruction.md                    | orchestration | Scanned            | 5     | 0.4       | 105    | -         |
| README.md                         | other         | Scanned            | 4     | 0.4       | 105    | -         |
| docs/guide.md                     | other         | Scanned            | 4     | 0.4       | 102    | -         |
| workflow.md                       | orchestration | Scanned            | 5     | 0.4       | 97     | -         |
| config/settings.yaml              | other         | Scanned            | 5     | 0.1       | 15     | -         |
| docs/architecture.mmd             | diagram       | Scanned            | 3     | 0.0       | 13     | -         |
| package.json                      | other         | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-1/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-2/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-3/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-4/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-5/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-6/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-7/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-8/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-9/index.js           | source        | Scanned            | 1     | 0.0       | 10     | -         |
| services/svc-1/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-2/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-3/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-4/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-5/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-6/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-7/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-8/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| services/svc-9/package.json       | other         | Scanned            | 1     | 0.0       | 4      | -         |
| assets/blob.json                  | other         | Binary             | 0     | 0.0       | 0      | -         |
| data-dump.json                    | other         | SkippedOversized   | 0     | 6766.5    | 0      | -         |
| assets/asset-0.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-1.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-10.png.txt.bak       | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-11.png.txt.bak       | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-2.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-3.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-4.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-5.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-6.png.txt.bak        | other         | SkippedUnsupported | 0     | 0.0       | 0      | -         |



| Path                       | Category | Status             | Lines | Size (KB) | Tokens | Dup group |
|----------------------------|----------|--------------------|-------|-----------|--------|-----------|
| assets/asset-7.png.txt.bak | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-8.png.txt.bak | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/asset-9.png.txt.bak | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/logo.svg            | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| assets/notes.docx          | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-1/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-2/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-3/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-4/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-5/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-6/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-7/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-8/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |
| services/svc-9/Dockerfile  | other    | SkippedUnsupported | 0     | 0.0       | 0      | -         |

## Skipped files and warnings

Skipped by reason: Binary = 1, SkippedOversized = 1, SkippedUnsupported = 23

| Path                        | Status             | Size (KB) | Reason                                                 |
|-----------------------------|--------------------|-----------|--------------------------------------------------------|
| assets/blob.json            | Binary             | 0.0       | File contains binary data                              |
| data-dump.json              | SkippedOversized   | 6766.5    | File is 6.6 MB which is above the 5 MB parse limit     |
| assets/asset-0.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-1.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-10.png.txt.bak | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-11.png.txt.bak | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-2.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-3.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-4.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-5.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-6.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-7.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-8.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |
| assets/asset-9.png.txt.bak  | SkippedUnsupported | 0.0       | Extension '.bak' is not in the supported analysis list |

| Path                      | Status             | Size (KB) | Reason                                                  |
|---------------------------|--------------------|-----------|---------------------------------------------------------|
| assets/logo.svg           | SkippedUnsupported | 0.0       | Extension '.svg' is not in the supported analysis list  |
| assets/notes.docx         | SkippedUnsupported | 0.0       | Extension '.docx' is not in the supported analysis list |
| services/svc-1/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-2/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-3/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-4/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-5/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-6/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-7/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-8/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |
| services/svc-9/Dockerfile | SkippedUnsupported | 0.0       | Extension 'none' is not in the supported analysis list  |