

Agentic Efficiency Report (Redacted)

agentic-crm · zip import · scan SCN-2026-01-01-0001

42

Wasteful + PartialScan
Weighted efficiency verdict

Redaction rules applied

File contents and code samples are removed. Exact file and folder names are replaced with path aliases such as dir-01/file-007.md. Counts, scores, category findings, penalties and savings estimates are preserved exactly as in the full report.

What we found

Overall efficiency score	42 / 100
Verdict	Wasteful (PartialScan)
Files discovered	94
Files parsed	69
Files skipped	25
Tokens analysed	55,038
Estimated monthly token waste	17,534,970
Estimated monthly credit waste	175.35 credits
Estimated monthly dollar waste	\$98.19
Savings range (low / high)	\$78.55 - \$117.83
Duplicate clusters found	2
Oversized context files	3
Overlapping agent groups	1
Review stages inferred	10
Agent-like files detected	46

Category scores

Category	Score	Weight	Penalty	What this means
Redundancy	0/100	25%	-108	2 groups of near-identical files and 4 repeated instruction blocks were found.
Token bloat	33/100	25%	-67	3 context or memory files are over 8,000 tokens.
Review overhead	64/100	20%	-36	10 review or approval steps were inferred from your instructions.
Agent sprawl	61/100	20%	-39	46 agent-style files were found (20 agent roles, 13 skills).
Architecture inefficiency	83/100	10%	-17	9 service directories against 9 hand written source files.

Top 5 waste drivers, in plain language

1. The same instruction block is pasted into 3 files

One block of instructions has been copied into several agent files. Only one copy is needed. The rest is paid for on every run.

Redundancy · 6,268,000 tokens/month · 62.68 credits · \$35.10

2. Context file is very large: file-045.md

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

Token bloat · 2,535,200 tokens/month · 25.35 credits · \$14.20

3. 10 separate review or approval steps were found

Your instructions describe many checks before work is accepted. Each extra check means another pass over the same material, which costs time and money.

Review overhead · 2,045,100 tokens/month · 20.45 credits · \$11.45

4. 46 agent-style files is more than this repo needs

There are many separate agent, skill and prompt files. Each one adds instructions that have to be loaded and kept in sync. Fewer, clearer files cost less and break less often.

Agent sprawl · 1,632,000 tokens/month · 16.32 credits · \$9.14

5. Context file is very large: file-038.md

This memory or context file is bigger than a healthy working budget. Large context files are re-read on every run, so most of the cost is paid again and again.

Token bloat · 1,553,200 tokens/month · 15.53 credits · \$8.70

Savings assumptions and formulas

- estimated_tokens = ceil(character_count / 4)
- Report credits: 1 credit = 100,000 tokens.
- Vendor billing: \$1.00 = 100 vendor credits; 1 vendor credit = 2,500 input tokens or 500 output tokens (output costs 5x input).
- Derived rates: \$4.00 per 1M input tokens, \$20.00 per 1M output tokens.
- Waste is assumed to be 90% input tokens and 10% output tokens.
- Each agent asset is assumed to be loaded 200 times per month.
- Aggregate savings show a +/-20% range.
- Similarity uses normalised text so whitespace-only and case-only changes are ignored.
- Files under 200 characters are excluded from duplicate detection.
- Scoring runs in two tiers. Tier 1 is the seven specified penalty rules with their hard caps. Tier 2 is severity scaling: for each category we measure the share of the monthly agent context budget that the category wastes, and deduct proportionally up to that category's tier 2 cap, with a category that wastes 50% or more of the budget taking the full deduction. Without tier 2 the capped rules alone could never produce a score below about 72, so the Wasteful and Critical bands would be unreachable. Both tiers are itemised in the score ledger.
- Monthly agent context budget for this repository = 10,933,400 tokens (54,667 agent asset tokens x 200 runs per month).

Setting	Value
Tokens per report credit	100,000
Vendor credits per dollar	100.0
Input tokens per vendor credit	2,500
Output tokens per vendor credit	500
\$ per 1M input tokens	\$4.00
\$ per 1M output tokens	\$20.00
Agent runs per month	200
Output token share of waste	10%
Aggregate variance	+/-20%
Rates last refreshed	never
Rates source	Product owner supplied defaults (not yet refreshed)

How the score was calculated

Penalty rule	Hits	Points each	Cap	Applied	Category
Near-duplicate agent/context file pair (similarity >= 0.80)	20	5	32	-32	redundancy
Repeated instruction block >= 150 chars in 3+ files	4	4	20	-16	redundancy
Context or memory file over 8,000 tokens	3	6	24	-18	token_bloat
Agent-like files above 12	34	2	20	-20	agent_sprawl
Review/approval stages above 4	6	5	15	-15	review_overhead
Microservice mismatch (>= 8 service dirs and < 120 source files)	1	12	12	-12	architecture_inefficiency
Monolith mismatch (1 service and >= 10 overlapping agent roles)	0	10	10	-0	architecture_inefficiency
Severity scaling - share of the monthly agent context budget wasted by duplication	80	0	60	-60	redundancy
Severity scaling - share of the monthly agent context budget wasted by oversized files	41	0	60	-49	token_bloat
Severity scaling - share of the monthly agent context budget wasted on review loops	19	0	55	-21	review_overhead
Severity scaling - share of the monthly agent context budget wasted on extra agents	16	0	60	-19	agent_sprawl
Severity scaling - share of the monthly agent context budget lost to the architecture	5	0	50	-5	architecture_inefficiency

Findings (14 issues)

Issue	Severity	Category	Title	Impacted files	Tokens /mo	Credits/mo	\$/mo
ISS-2026-07-29-0006	critical	redundancy	The same instruction block is pasted into 3 files	dir-05/file-038.md, dir-05/file-039.md, dir-07/file-045.md	6,268,000	62.68	\$35.10
ISS-2026-07-29-0007	high	token_bloat	Context file is very large: file-045.md	dir-07/file-045.md	2,535,200	25.35	\$14.20
ISS-2026-07-29-0011	high	review_overhead	10 separate review or approval steps were found	dir-01/file-002.md, dir-01/file-003.md, dir-01/file-004.md, dir-01/file-005.md	2,045,100	20.45	\$11.45
ISS-2026-07-29-0010	high	agent_sprawl	46 agent-style files is more than this repo needs	dir-01/file-002.md, dir-01/file-003.md, dir-01/file-004.md, dir-01/file-005.md (+21 more)	1,632,000	16.32	\$9.14
ISS-2026-07-29-0008	high	token_bloat	Context file is very large: file-038.md	dir-05/file-038.md	1,553,200	15.53	\$8.70
ISS-2026-07-29-0001	high	redundancy	20 files are near copies of each other	dir-01/file-002.md, dir-01/file-003.md, dir-01/file-004.md, dir-01/file-005.md (+16 more)	1,055,200	10.55	\$5.91
ISS-2026-07-29-0003	medium	redundancy	The same instruction block is pasted into 43 files	dir-01/file-002.md, dir-01/file-003.md, dir-01/file-004.md, dir-01/file-005.md (+39 more)	764,400	7.64	\$4.28
ISS-2026-07-29-0012	high	architecture_inefficiency	Split into many services but there is little code in them	[hidden], [hidden], [hidden], [hidden] (+5 more)	546,670	5.47	\$3.06
ISS-2026-07-29-0009	medium	token_bloat	Context file is very large: file-039.md	dir-05/file-039.md	401,000	4.01	\$2.25
ISS-2026-07-29-0004	medium	redundancy	The same instruction block is pasted into 23 files	dir-01/file-002.md, dir-01/file-003.md, dir-01/file-004.md, dir-01/file-005.md (+19 more)	391,600	3.92	\$2.19
ISS-2026-07-29-0005	low	redundancy	The same instruction block is pasted into 20 files	dir-01/file-002.md, dir-01/file-003.md, dir-01/file-004.md, dir-01/file-005.md (+16 more)	182,400	1.82	\$1.02
ISS-2026-07-29-0013	low	agent_sprawl	13 skill files, and some repeat each other	dir-18/file-081.md, dir-18/file-082.md, dir-18/file-083.md, dir-18/file-084.md (+9 more)	72,000	0.72	\$0.40
ISS-2026-07-29-0014	low	review_overhead	10 orchestration files create extra hand-offs	file-043.md, file-044.md, file-046.md, dir-08/file-048.md (+6 more)	50,000	0.50	\$0.28
ISS-2026-07-29-0002	low	redundancy	2 files are near copies of each other	file-044.md, file-046.md	38,200	0.38	\$0.21

Recommended actions, ranked by impact then effort

#	Action	Category	Impact	Effort	Tokens/mo saved	\$/mo saved
1	Move the shared block into one file and reference it from the others.	Redundancy	High	Small	764,400	\$4.28
2	Keep at most four review gates. Combine the overlapping ones into a single checklist.	Review overhead	High	Medium	2,045,100	\$11.45

#	Action	Category	Impact	Effort	Tokens/mo saved	\$/mo saved
3	Keep the most complete file, delete or shrink the copies, and link to the kept file instead.	Redundancy	High	Medium	1,055,200	\$5.91
4	Group agents by job. Merge the ones that overlap and delete the ones nobody calls.	Agent sprawl	High	Large	1,632,000	\$9.14
5	Fold the thin services back together until each one has a clear, separate job.	Architecture inefficiency	High	Large	546,670	\$3.06
6	Keep one skill file per capability and delete the near copies.	Agent sprawl	Medium	Small	72,000	\$0.40
7	Split this file into a short always-on summary plus detail files that load only when needed.	Token bloat	Medium	Medium	2,535,200	\$14.20
8	Describe the flow once in a single orchestrator file and let the agents stay simple.	Review overhead	Medium	Medium	50,000	\$0.28

File inventory by category

Group	Files	Parsed	Skipped	Tokens
Agents	20	20	0	5,564
Skills	13	13	0	1,403
Context and memory	3	3	0	46,447
Prompt and orchestration	10	10	0	1,253
Source code	9	9	0	90
Diagrams	1	1	0	13
Docs	2	2	0	207
Other text assets	36	11	25	61

Largest 60 files by estimated tokens

Path	Category	Status	Lines	Size (KB)	Tokens	Dup group
dir-07/file-045.md	context_memory	Scanned	8	80.8	20,676	-
dir-05/file-038.md	context_memory	Scanned	7	61.6	15,766	-
dir-05/file-039.md	context_memory	Scanned	5	39.1	10,005	-
dir-02/file-007.md	agent	Scanned	23	1.1	288	DUP-001
dir-01/file-006.md	agent	Scanned	23	1.1	280	DUP-001
dir-02/file-008.md	agent	Scanned	23	1.1	280	DUP-001
dir-02/file-009.md	agent	Scanned	23	1.1	280	DUP-001
dir-01/file-003.md	agent	Scanned	23	1.1	279	DUP-001
dir-02/file-010.md	agent	Scanned	23	1.1	279	DUP-001
dir-01/file-013.md	agent	Scanned	23	1.1	279	DUP-001
dir-01/file-004.md	agent	Scanned	23	1.1	278	DUP-001
dir-01/file-015.md	agent	Scanned	23	1.1	278	DUP-001
dir-01/file-018.md	agent	Scanned	23	1.1	278	DUP-001
dir-01/file-002.md	agent	Scanned	23	1.1	277	DUP-001
dir-01/file-005.md	agent	Scanned	23	1.1	277	DUP-001
dir-01/file-012.md	agent	Scanned	23	1.1	277	DUP-001
dir-01/file-017.md	agent	Scanned	23	1.1	277	DUP-001
dir-01/file-019.md	agent	Scanned	23	1.1	277	DUP-001
dir-01/file-020.md	agent	Scanned	23	1.1	277	DUP-001
dir-01/file-011.md	agent	Scanned	23	1.1	276	DUP-001
dir-01/file-014.md	agent	Scanned	23	1.1	276	DUP-001
dir-01/file-021.md	agent	Scanned	23	1.1	276	DUP-001
dir-01/file-016.md	agent	Scanned	23	1.1	275	DUP-001
file-046.md	orchestration	Scanned	15	0.8	206	DUP-002
file-044.md	orchestration	Scanned	7	0.7	191	DUP-002
dir-08/file-048.md	orchestration	Scanned	5	0.4	109	-
dir-08/file-049.md	orchestration	Scanned	5	0.4	109	-
dir-08/file-050.md	orchestration	Scanned	5	0.4	109	-
dir-08/file-051.md	orchestration	Scanned	5	0.4	109	-
dir-08/file-052.md	orchestration	Scanned	5	0.4	109	-
dir-08/file-053.md	orchestration	Scanned	5	0.4	109	-
dir-18/file-088.md	skill	Scanned	5	0.4	109	-
dir-18/file-091.md	skill	Scanned	5	0.4	109	-
dir-18/file-092.md	skill	Scanned	5	0.4	109	-
dir-18/file-081.md	skill	Scanned	5	0.4	108	-
dir-18/file-082.md	skill	Scanned	5	0.4	108	-
dir-18/file-083.md	skill	Scanned	5	0.4	108	-

Path	Category	Status	Lines	Size (KB)	Tokens	Dup group
dir-18/file-084.md	skill	Scanned	5	0.4	108	-
dir-18/file-089.md	skill	Scanned	5	0.4	108	-
dir-18/file-093.md	skill	Scanned	5	0.4	108	-
dir-18/file-085.md	skill	Scanned	5	0.4	107	-
dir-19/file-086.md	skill	Scanned	5	0.4	107	-
dir-18/file-087.md	skill	Scanned	5	0.4	107	-
dir-18/file-090.md	skill	Scanned	5	0.4	107	-
file-043.md	orchestration	Scanned	5	0.4	105	-
file-001.md	other	Scanned	4	0.4	105	-
dir-06/file-042.md	other	Scanned	4	0.4	102	-
file-094.md	orchestration	Scanned	5	0.4	97	-
dir-04/file-037.yaml	other	Scanned	5	0.1	15	-
dir-06/file-041.mmd	diagram	Scanned	3	0.0	13	-
file-047.json	other	Scanned	1	0.0	10	-
dir-09/file-055.js	source	Scanned	1	0.0	10	-
dir-10/file-058.js	source	Scanned	1	0.0	10	-
dir-11/file-061.js	source	Scanned	1	0.0	10	-
dir-12/file-064.js	source	Scanned	1	0.0	10	-
dir-13/file-067.js	source	Scanned	1	0.0	10	-
dir-14/file-070.js	source	Scanned	1	0.0	10	-
dir-15/file-073.js	source	Scanned	1	0.0	10	-
dir-16/file-076.js	source	Scanned	1	0.0	10	-
dir-17/file-079.js	source	Scanned	1	0.0	10	-

Skipped files and warnings

Skipped by reason: Binary = 1, SkippedOversized = 1, SkippedUnsupported = 23

Path	Status	Size (KB)	Reason
dir-03/file-034.json	Binary	0.0	File contains binary data
file-040.json	SkippedOversized	6766.5	File is 6.6 MB which is above the 5 MB parse limit
dir-03/file-022.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-023.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-024.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-025.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-026.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-027.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-028.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-029.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-030.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-031.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-032.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list
dir-03/file-033.bak	SkippedUnsupported	0.0	Extension '.bak' is not in the supported analysis list

Path	Status	Size (KB)	Reason
dir-03/file-035.svg	SkippedUnsupported	0.0	Extension '.svg' is not in the supported analysis list
dir-03/file-036.docx	SkippedUnsupported	0.0	Extension '.docx' is not in the supported analysis list
dir-09/file-054	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-10/file-057	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-11/file-060	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-12/file-063	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-13/file-066	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-14/file-069	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-15/file-072	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-16/file-075	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list
dir-17/file-078	SkippedUnsupported	0.0	Extension 'none' is not in the supported analysis list

Draft replacement files

2 draft replacement files were generated. Draft contents are removed from the redacted report. Aggregate reduction is preserved below.

Alias	Type	Impact	Effort	Original tokens	Draft tokens	Reduction
dir-07/file-045.md	memory	High	Large	20,676	872	-95.8%
dir-05/file-038.md	context	High	Medium	15,766	806	-94.9%